

Numerical Mathematics II for Engineers Tutorial 0

Topics : Reminder on Python and NumPy, Finite difference quotients, solving the 1D Poisson equation with finite differences.

Discussion in the tutorials of the week 20–26 October 2025

Disclaimer: This exercise sheet looks very lengthly, but the exercises are only a revision of things that should already be known; they should therefore not take too long to complete.

Exercise 0.1: Introduction to Numpy and Matplotlib

This exercise is an introduction to the NumPy package, inspired from the official quickstart guide¹. NumPy is a package for the multidimensional array object `ndarray`. It allows fast operations on large arrays, and has become the basis for scientific computing in Python. Like classical Python `list`, NumPy `array` are tables of elements, but they have the following characteristics:

- arrays have a fixed, predetermined size,
- all elements of an array must have the same type,
- dimensions are called "axes".

NumPy enables **vectorization**: explicit looping is generally avoided in the code. Example: assume two arrays u, v of dimension $(N, 1)$ are given (i.e., essentially vectors). We want to know their sum. In traditional programming languages, one would do that by using loops

```
1 w = [0 for k in range(N)]
2 for k in range(N):
3     w[k] = u[k] + v[k]
```

In a vectorized form, this would just be

```
1 w = u + v
```

which is faster. It also makes it easier to understand what the code is doing.

Vectorization also involves the usage of built-in functions when possible. As a second example, we compare the runtime to compute the scalar product of two vectors u, v of dimension N , either using loops or using vectorization and built-in NumPy functions.

```
1 ## Using loops
2 for i in range(N):
3     w[i] = u[i]*v[i]
4
5 ## Using the built-in NumPy functions
6 w = np.dot(u,v)
```

¹<https://numpy.org/doc/stable/user/quickstart.html>

N	Time with loop (s)	Time with NumPy (s)
10	3.34e-06	1.74e-05
1000	6.29e-05	5.48e-06
1e6	7.21e-02	7.16e-03
5e6	3.51e-01	1.82e-02
1e7	6.91e-01	3.75e-02

Table 1: Runtime for computing the scalar product of two vectors of size N

The runtimes obtained for increasing values of N are presented in Table 1. For small N , the times are comparable, but for large N , the NumPy implementation is at least ten times faster.

Built-in and vectorized operations are optimized and pre-compiled in C, making them very fast and adapted for scientific computing. **Whenever possible, use vectorized code!**

Now to the exercise: install numpy (preferably in a virtual environment), open an interactive console and do the following tasks²:

- (a) Run `import numpy as np`.
- (b) Define `a,b,c` as follows:

```

1 a = np.array([1, 0, 0])
2 b = np.array([[1., 0., 0.],[0., 1., 2.]])
3 c = np.array([[[0, 1, 2, 3], [4, 5, 6, 7], [8, 9, 10, 11]],
4                 [[12, 13, 14, 15], [16, 17, 18, 19],
5                  [20, 21, 22, 23]]])

```

For each array, check the following properties: `ndim`, `nshape`, `size`, `dtype`. Example: `a.ndim`

- (c) Arrays can be created in various ways. In question 1.1(a), you converted lists into arrays, but the built-in functions are more suited. Try the following functions:

```

1 np.zeros((3, 4))
2 np.ones((2, 3, 4))
3 np.arange(10, 30, 5)
4 np.linspace(0, 2, 9)

```

Once an array is created, its shape can be changed. Try the following operations:

```

1 a = np.ones((3, 4))
2 a.flatten()
3 a.reshape(6, 2)
4 a.T

```

Note that those operations return a new array, they do not change `a`. Now, use `arange` and `reshape` to create the array `c` from Question 1.1(b) in one line of code.

- (d) Most operations on arrays apply elementwise. Try the following manipulations:

```

1 a = np.array([20, 30, 40, 50])
2 b = np.arange(4)
3 c = a - b

```

²It is unfortunately not possible to provide support from our side for the installation of Python, NumPy and Matplotlib. Using tools like ChatGPT though this should hopefully be straightforward.

```

4 b**2
5 a < 35
6 a *= 3
7
8 ## Classical functions are available in NumPy
9 10 * np.sin(a)
10 np.exp(B)
11 np.sqrt(B)

```

Finally, note the difference between matrix dot product and elementwise product:

```

1 A = np.array([[1, 1],
2                 [0, 1]])
3 B = np.array([[2, 0],
4                 [3, 4]])
5 A * B      # elementwise product
6 A @ B      # matrix product
7 A.dot(B)   # another matrix product

```

- (e) Like Python lists, arrays can be indexed, sliced and iterated over. Indices for different axes are separated by a comma. Try the following operations:

```

1 def f(x, y):
2     return 10 * x + y
3 b = np.fromfunction(f, (5, 4), dtype=int)
4 b[2, 3]
5 b[0:5, 1]    # each row in the second column of b
6 b[:, 1]       # equivalent to the previous example
7 b[1:3, :]    # each column in the second and third row of b

```

- (f) To visualize our results, we will use the `matplotlib` package³. Install matplotlib (again, preferably in a virtual environment) and run the following code:

```

1 import matplotlib.pyplot as plt
2 x = np.linspace(0, 2*np.pi)
3 plt.plot(x, np.sin(x), "x-")
4 plt.xlabel("$x$")           # LaTeX can be used for plotting
5 plt.ylabel("$\sin(x)$")    # But be careful with the "\ command
6 plt.show()

```

There are many options to customize plots, you will see some examples in the templates of this homework.

For more information on NumPy and Matplotlib, look at their officiel documentation: <https://numpy.org/doc/stable/index.html> for NumPy and <https://matplotlib.org/stable/> for Matplotlib.

Exercise 0.2: Revision on Landau notation

This exercise is a review on the Landau, or "big O", notation.

- (a) The following polynomial is given:

$$f(x) = 2x^2 + 3x^3 + 8x^5.$$

Give the behavior of $f(x) = \mathcal{O}(x^?)$ for the limit of $x \rightarrow 0$ and for the limit of $x \rightarrow \infty$.

³<https://matplotlib.org/>

(b) The following polynomial is given:

$$f(x) = 2x^2 + 3x^3 + 8x^5.$$

Give the behavior of $f(x) = 2x^2 + \mathcal{O}(x^7)$ for the limit of $x \rightarrow 0$.

Exercise 0.3: Revision on convergence curves and tables

This exercise is a reminder on convergence orders, and on how to compute and display them. For the programming questions, the template `convergence_template.py` is provided. As a model problem, we will consider two different quadrature rules to integrate a function $f \in C^4([a, b])$ over the interval $[a, b] \subset \mathbb{R}$, i.e., to approximately compute

$$I = \int_a^b f(x) dx.$$

For the numerical integration, the interval $[a, b]$ is discretized into n subintervals $[x_j, x_{j+1}]$, with $x_j = a + jh$ and $h = (b - a)/n$. We will compare the following quadrature rules:

- the **composite Trapezoidal rule** (on each interval, the two end points are used)

$$\int_a^b f(x) dx \approx T(h) = h \left(\frac{f(a) + f(b)}{2} + \sum_{j=1}^{n-1} f(x_j) \right),$$

- the **composite Simpson's rule** for an even number of subdivision n (on each interval, the two end points and the midpoint are used)

$$\int_a^b f(x) dx \approx S(h) = \frac{h}{3} \left(f(a) + 4 \sum_{j=1}^{n/2} f(x_{2i-1}) + 2 \sum_{j=1}^{n/2-1} f(x_{2i}) + f(b) \right).$$

If the function f is sufficiently smooth (which we assume for this exercise), then there holds for the error $E(h)$:

- for the Trapezoidal rule: $E_T(h) = |I - T(h)| = \mathcal{O}(h^2)$ for $h \rightarrow 0$;
- for the Simpson's rule $E_S(h) = |I - S(h)| = \mathcal{O}(h^4)$ for $h \rightarrow 0$.

That means, if we are already in the asymptotic region, i.e., if h is fine enough, then if we reduce the interval length by a factor of 2, i.e., $h \rightarrow \frac{h}{2}$, we expect that

- for the Trapezoidal rule: the error is reduced by a factor $2^2 = 4$,
- for Simpson's rule: the error is reduced by a factor of $2^4 = 16$.

We integrate the function $f(x) = 2 + \cos(x)$ over the interval $[0, \pi/2]$.

- Compute the exact integral.
- Implement the Trapezoidal rule in the function `approximate_integral` in the file `convergence_template.py`.

- (c) For the Simpson's rule, we use the `simpson` function from the `scipy.integrate` package. Compute the relative error between the approximated integral and the exact integral for $n \in \{2, 4, 8, 16, 32\}$ for both quadrature rules. Plot both errors in the same picture using `plot` from the `matplotlib.pyplot` package. Which rule is more accurate? Can you easily read the convergence order on the curve?
- (d) To better distinguish between both curves, plot the relative errors in **logarithmic scaling** using `loglog` from the `matplotlib.pyplot` package. Also plot the curves $f(h) = h^2$ and $f(h) = h^4$.
 You should get straight lines for the curves h^2 and h^4 , explain why. Which convergence order do you observe for the Trapezoidal and for the Simpson's rule?
- (e) Using `loglog` plots is very helpful for reading off convergence orders. Another very good alternative is to use error tables to compute the orders. Assuming that we have convergence of the form h^p and we want to determine p , we can use

$$p \approx \log_2 \left(\frac{E(2h)}{E(h)} \right). \quad (1)$$

A derivation of the formula can be found in the appendix A.1. In the code, compute the table of order from the errors. Since you computed the errors for five different h , you should get four values for the order, for each quadrature rule.

Exercise 0.4: Finite difference quotients

We first recall the definitions and results of the lecture. We consider an open interval $I \subset \mathbb{R}$ and a function $f \in C^4(\bar{I})$. For all $x \in I$ and $h \in \mathbb{R}$ such that $x + h \in I$, the forward difference quotient is defined by

$$D^+ f(x) = \frac{f(x+h) - f(x)}{h} \quad (\text{FD}).$$

In the lecture we have proved, using Taylor expansion, that there holds

$$D^+ f(x) = f'(x) + hC_1, \quad (2)$$

with

$$|C_1| \leq \frac{1}{2} \max_{\xi \in [x, x+h]} |f'(\xi)|. \quad (3)$$

The equations (2)-(3) imply that

$$\left| f'(x) - D^+ f(x) \right| = \mathcal{O}(h) \quad (h \rightarrow 0),$$

hence, we say that (FD) **converges with first order / is an approximation of first order** to $f'(x)$.

- (a) For all $x \in I$ and $h \in \mathbb{R}$ such that $x + h \in I$ and $x - h \in I$, the central difference quotient is defined by

$$D_1^c f(x) = \frac{f(x+h) - f(x-h)}{2h} \quad (\text{CD}).$$

Prove the equation from the lecture

$$D_1^c f(x) = f'(x) + h^2 C_2, \quad (4)$$

with

$$|C_2| \leq \frac{1}{6} \max_{\xi \in [x-h, x+h]} |f^{(3)}(\xi)|. \quad (5)$$

We say that (CD) **converges with second order / is an approximation of second order** to $f'(x)$.

- (b) Implement the schemes (FD) and (CD) in the template `FD_template.py`. To check your implementation, apply the scheme (FD) to a linear function, and the scheme (CD) to a quadratic function: in each case, the derivative should be approximated to machine precision.
- (c) We choose $f(x) = \exp(x)$ and $x = 1$, and want to test **numerically** the convergence orders of (FD) and (CD). Plot the relative error for (FD) and (CD) in a `loglog` plot (with h on the x -axis and the relative error on the y -axis) for $h = 10^{-k}$ with $k \in \{1, 2, 3, 4, 5\}$. Also plot the functions h and h^2 for comparison. Does the result meet your expectation?
- (d) Now do the same for $h = 10^{-k}$ with $k = 1, 2, \dots, 11, 12$. What do you observe now? Can you explain the phenomenon?

A Appendix

A.1 Computing the convergence order

Derivation of the formula (1): assume that the error has the form

$$E(h) = Ch^p + \mathcal{O}(h^{p+1})$$

for some constant C independent of h .

We investigate how the error changes when we divide the step length h by 2. Assuming that the constant C does not change, we get

$$E(2h) = C(2h)^p + \mathcal{O}((2h)^{p+1}) = 2^p Ch^p + \mathcal{O}(h^{p+1})$$

and then

$$\frac{E(2h)}{E(h)} = \frac{2^p Ch^p + \mathcal{O}(h^{p+1})}{Ch^p + \mathcal{O}(h^{p+1})} = \frac{2^p + \mathcal{O}(h)}{1 + \mathcal{O}(h)} = 2^p + \mathcal{O}(h).$$

Taking the logarithm on both sides, we get

$$p \approx \log_2 \left(\frac{E(2h)}{E(h)} \right).$$